

Starting from Information

Ruichen Wang
wangrc@2345.com

November 20, 2018

Abstract

Starting with the introducing the origins of information, extend to entropy and its families. and some introduction and explanation of entropy related algorithms, like log-likelihood, softmax classification. Then we will talk about one basic algorithm solving intractable problems, which is called variational bayesian inference, and it's closely related algorithms like latent dirichlet allocation(LDA), variational auto encoder(VAE), generative adversarial network(GAN),etc.

Contents

1	What is Information?	2
1.1	Defination of Information	2
1.2	Property of Information	2
2	Entorpy (Expectation of Information)	3
2.1	Shannon's Information Theory	3
2.2	Property of Entropy	3
3	Families of Entropy	3
3.1	Cross-Entropy	3
3.1.1	Relation to Log-likelihood	3
3.1.2	Cross-entropy Loss in Classification	4
3.1.3	Relationship with Softmax	4
3.2	Kullback-Leibler Divergence	5
3.2.1	Interpretations	5
3.2.2	Property of KL	5
3.2.3	Applications	5
4	Variational Bayesian Inference	6
4.1	Variational Inference	6
4.2	KL Divergence Measure	6
4.3	Evidence Lower Bound	7
4.3.1	Jensen's Inequality	7

4.3.2	KL Transformation	7
4.3.3	Relationship with EM	8
4.4	Mean Field Theory	8
4.4.1	Mean Field Approximation	9
4.4.2	Mean Field Method	9
4.5	Coordinate Ascent Variational Inference	9
5	Latent Dirichlet Allocation	10
5.1	Conjugate Distributions	11
5.1.1	Beta-Binomial Distribution	11
5.1.2	Dirichlet-Multinomial Distribution	12
5.2	Steps of LDA	12
5.3	Variational Inference for LDA	13

1 What is Information?

1.1 Defination of Information

How to measure the information of certain event in a mathematics?

Given x is certain event, $P(x)$ is probability which event x happens. Intuitively, the information should have inverse proportion to the probability, higher probability should have lower information, which is

$$I(x) = \frac{1}{P(x)}$$

As $p(x) \in [0, 1]$, we also want the information more stable, and remove the division for calculation convience, so we can re-define it as:

$$I(x) = \log \frac{1}{P(x)} = -\log P(x)$$

1.2 Property of Information

As a result, the $-\log P(x)$ has every properties we want:

- Lower probability, higher information
- Higher probability, lower information
- Multi-event happens, the probability is multiplied. the information is summed

$$P(x_1, x_2) = P(x_1) * P(x_2)$$

$$\log P(x_1, x_2) = \log P(x_1) + \log P(x_2)$$

$P(x) \in [0, 1]$, and larger $P(x)$ has smaller information.

2 Entorpy (Expectation of Information)

2.1 Shannon's Information Theory

Claude Elwood Shannon(1916-2001).

1937 MIT Master degree.

1940 MIT Ph.D degree from MIT.

1948 Published a landmark paper 'A mathematical Theory of Communication'.

Entropy is defined as the expectation of information certain event carries:

$$H(x) = E[I(x)] = \sum_{i=1}^n P(x_i)I(x_i) = - \sum_{i=1}^n P(x_i)\log P(x_i)$$

2.2 Property of Entropy

The property of the entropy is quite simple

- Higher probability, the less information, the lower entropy
- Non-negative, every event has some information
- Cumulative, multile events happens, the information is the sum of them.

3 Families of Entropy

3.1 Cross-Entropy

It is often the case that we don't know $P(x)$ yet, so we make an 'artificial' probability distribution $Q(x)$. How can we measure the cost as we using $Q(x)$ to approximate $P(x)$? We define corss entropy as:

$$H(P, Q) = - \sum_x P(x)\log Q(x)$$

In practice Given a test set N observed, which comes from a Monte Carlo sampling of the true distribution $P(x)$. Cross entropy is calculated using :

$$H(T, Q) = - \frac{1}{N} \sum_{i=1}^N \log Q(x_i)$$

3.1.1 Relation to Log-likelihood

for the maximum likelihood estimation (MAE), we want:

$$\arg \max \prod_i q_i^{N_{p_i}}$$

So log-likelihood, divided by N is :

$$\frac{1}{N} \log \prod_i q_i^{N p_i} = \sum_i p_i \log q_i = -H(p, q)$$

So maximum the likelihood is the same as minimizing the cross entropy.

3.1.2 Cross-entropy Loss in Classification

In machine learning, cross-entropy loss is widely used, it often defines as :

$$L = -y \log(y') = H(y, y')$$

It describes the distance between the prediction and truth.

3.1.3 Relationship with Softmax

As the softmax probability and cross-entropy loss is so so common, and they often work together. But why? Because the simplicity of the derivative. Softmax function:

$$p_i = \frac{e^{a_i}}{\sum_{k=1}^N e^{a_k}}$$

Derivative of softmax $\frac{\partial p_i}{\partial \alpha_j}$:

$$\frac{\partial p_i}{\partial \alpha_j} = \begin{cases} p_i(1 - p_j) & i = j \\ -p_j * p_i & i \neq j \end{cases}$$

The cross entropy loss:

$$L = - \sum_i y_i \log p_i$$

Derivative of cross entropy loss:

$$\frac{\partial L}{\partial o_i} = - \sum y_k \frac{1}{p_k} * \frac{\partial p_k}{\partial o_i}$$

From the dervative of softmax we derived earlier,

$$\begin{aligned} \frac{\partial L}{\partial o_i} &= -y_i(1 - p_i) - \sum_{k \neq i} y_k \frac{1}{p_k} (-p_k * p_i) \\ &= p_i(y_i + \sum_{k \neq i} y_k) - y_i = p_i - y_i \end{aligned}$$

This is why we often use softmax and cross entropy together, The gradient is quite simple and elegant to calculate.

3.2 Kullback-Leibler Divergence

KL divergence is also called relative entropy. It is a measure of how one probability distribution is different from a second.

For discrete probability distributions P and Q defined on the same probability space, the KL divergence from Q to P (Q with respect to P) is defined as :

$$\begin{aligned} D_{KL}(P \parallel Q) &= H(P, Q) - H(P) \\ &= - \sum_i P(i) \log\left(\frac{Q(i)}{P(i)}\right) \end{aligned}$$

Which means the more entropy using Q generates with respect to original distribution P .

3.2.1 Interpretations

In machine learning, $D_{KL}(P \parallel Q)$ is often called the information gain achieved if Q is used instead of P .

Expressed in the language of Bayesian inference, $D_{KL}(P \parallel Q)$ is a measure of the information gained when one revises one's beliefs from the prior probability distribution Q to the posterior probability distribution P .

In applications, P typically represents the true distribution of data. Q represents the model. Minimize $D_{KL}(P \parallel Q)$ can be a good solution to find a Q that closest to P .

3.2.2 Property of KL

- Non-negative

As a result known as Gibbs's inequality, with $D_{KL}(P \parallel Q)$ zero if and only if $P = Q$.

- Asymmetric

$$D_{KL}(P \parallel Q) \neq D_{KL}(Q \parallel P)$$

we can define symmetrised divergence as:

$$\frac{D_{KL}(P \parallel Q) + D_{KL}(Q \parallel P)}{2}$$

3.2.3 Applications

It is widely used in generative models. You can find it in NLP, computer vision, robotics, biology. We will use this in next section. Here I just put the bayesian equation here. :)

$$\begin{aligned} \text{posterior} &= \frac{\text{likelihood} * \text{prior}}{\text{evidence}} \\ p(z|x) &= \frac{p(x|z)p(z)}{p(x)} \end{aligned}$$

4 Variational Bayesian Inference

4.1 Variational Inference

The idea behind variational inference is to first posit a family of densities and then to find the member of that family which is close to the target. Closeness is measured by **Kullback-Leibler divergence**.

Variational inference is widely used to approximate posterior densities for Bayesian models, an alternative strategy to Markov chain Monte Carlo (MCMC) sampling. MCMC algorithms try to sample a Markov chain, while the variational algorithms solve an optimization problem. Variational inference tends to be faster and easier to scale to large data and high dimension.

Variational inference has a close relationship with EM algorithm. You can view VAE, GAN as a certain form of variational inference. Variational inference doesn't have the global optimal point, which makes VAE/GAN very hard to train or converge. That's why you can find a lot of papers introducing how to train VAE/GAN more stably.

***Notes** Compare MCMC with Metropolis-Hasting (MH). MH larger the acceptance ratio α . When extended to high dimensions, it is called Gibbs sampling. Actually, they are based on the same idea - **Bayesian stationary distribution**.

Question description Suppose we have observations x , and hidden variables z , and some fixed parameters α . What we want is the posterior distribution.

$$p(z|x, \alpha) = \frac{p(z, x|\alpha)}{\int_z p(z, x|\alpha) dz}$$

In many cases, the $\int_z p(z, x|\alpha) dz$ is intractable. We don't know how to compute it especially in high dimensions.

Solution The main idea behind variational methods is to pick a family of distributions over the latent variables with its own **variational parameters**.

$$q(z_{1:m}|v)$$

Then find v to make q close to the posterior.

4.2 KL Divergence Measure

As mentioned above, we can use KL for this variational inference:

$$D_{KL}(q(z) \parallel p(z|x)) = E_{q(z)} \left[\log \frac{q(z)}{p(z|x)} \right]$$

Intuitively, According to this formula, there are three cases:

- If q is low, then we don't care (Because of the expectation)

- If q is high and p is high, good :)
- If q is high and p is low, bad :(

4.3 Evidence Lower Bound

Actually we can not minimize KL divergence. But we can minimize another function which is equal to this. This is evidence lower bound (ELBO).

4.3.1 Jensen's Inequality

Jensen's inequality are widely used in EM algorithm. In convex function, we have :

$$f(tx_1 + (1-t)x_2) \leq tf(x_1) + (1-t)f(x_2)$$

In the context of probability theory, if X is a random variable, and φ is a convex function, then:

$$\varphi(E[x]) \leq E[\varphi(x)]$$

Back to the problem, we have observations x^1, x^2, \dots, x^n , we want $p(x^i)$ get the max probability. Using MLE on it, which is the sum of the log-likelihood,

$$\log p_\theta(x^1, x^2, \dots, x^n) = \sum_{i=1}^N \log p_\theta(x^i)$$

and

$$\begin{aligned} \log p(x) &= \log \int_z p(x, z) \\ &= \log \int_z p(x, z) \frac{q(z)}{q(z)} \\ &= \log \left(E_q \left[\frac{p(x, z)}{q(z)} \right] \right) \\ &\geq E_q \left[\log \frac{p(x, z)}{q(z)} \right] \\ &\geq E_q[\log p(x, z)] - E_q[\log q(z)] \end{aligned}$$

Note the second term is the entropy

But what does this have to do with the KL divergence?

4.3.2 KL Transformation

As mentioned, we want $q(z)$ and $q(z|x)$ are close to each other:

$$\begin{aligned} KL(q(z)||p(z|x)) &= E_q \left[\log \frac{q(z)}{p(z|x)} \right] \\ &= E_q[\log q(z)] - E_q[\log p(z|x)] \end{aligned}$$

As we know,

$$p(z|x) = \frac{p(z, x)}{p(x)}$$

so

$$\begin{aligned} KL(q(z)||p(z|x)) &= E_q[\log q(z)] - E_q[\log p(z, x)] + E_q[\log p(x)] \\ &= -(E_q[\log p(z, x)] - E_q[\log q(z)]) + \log p(x) \end{aligned}$$

The first term is ELBO we just met.

The formula can also be written as :

$$\log p(x) = KL(q(z)||p(z|x)) + (E_q[\log p(z, x)] - E_q[\log q(z)])$$

As I mentioned before, For two different distributions, KL divergence is always non-negative. and $p(x)$ is the observation evidence, which is fixed. So minimizing the KL divergence is the same as maximizing the ELBO. This is also called as the variational lower bound.

4.3.3 Relationship with EM

EM algorithm is also known as a famous method to find the distributions of latent variables. Unlike variational inference we are going to talk about, EM algorithm use the fact that ELBO is equal to the $p(x)$ when $q(z) = p(z|x)$. EM **alternates** between computing $p(z|x)$ (E step), and optimizing it with respect to the model parameters(M step). The biggest difference is EM assume $p(z|x)$ is computable and fix the parameter, use it, while variational inference use bayesian setting and apply to the models we can not compute.

* EM is out the scope of this article, I don't want to go into too detail about it. Actually the formula below explains pretty clear.

E-step:

$$q(z) := p(z|x; \theta)$$

M-step:

$$\theta := \arg \max_{\theta} -KL(q(z)||p(z|x; \theta))$$

4.4 Mean Field Theory

Mean field theory is also called **self-consistent field theory**. It studies the behavior of large and complex stochastic models by studying a simpler model. Such models consider a large number of small individual components that interact with each other.

The effect of all the other individuals on any given individual is approximated by a single averaged effect, thus reducing a many-body problem to a one-body problem.

4.4.1 Mean Field Approximation

We assume each variable is independent. Using this theory, we can write:

$$q(z_{1:m}) = \prod_{i=1}^m q(z_i)$$

$$E_q[\log q(z_{1:m})] = \sum_{j=1}^m E_{q_j}[\log q(z_j)]$$

Also we have the chain rule

$$p(z_{1:m}, x_{1:n}) = p(x_{1:n}) \prod_{j=1}^m p(z_j | z_{1:(j-1)}, x_{1:n})$$

4.4.2 Mean Field Method

Note that the order of j is irrelevant. Based on this theory, we can rewrite the lower bound as:

$$\mathcal{L} = \sum_{j=1}^m E_q[\log p(z_j | z_{1:(j-1)}, x_{1:n})] - E_{q_j}[\log q(z_j)]$$

Consider the variable z_j comes last:

$$\mathcal{L} = \log p(x_{1:n}) + E_q[\log p(z_j | z_{-j}, x)] - E_{q_j}[\log q(z_j)]$$

And we can remove the first term because it's irrelevant to $q(z_j)$, the \mathcal{L} can be written as

$$\begin{aligned} \arg \min_{q_j} \mathcal{L} &= E_q[\log p(z_j | z_{-j}, x)] - E_{q_j}[\log q(z_j)] \\ &= \int q(z_j) E_{-j}[\log p(z_j | z_{-j}, x)] dz_j - \int q(z_j) \log q(z_j) dz_j \end{aligned}$$

4.5 Coordinate Ascent Variational Inference

Let's treat $q(z_j)$ as $f(x)$. For simplicity, I convert the formula above into this:

$$\frac{d\mathcal{L}}{dq(z_j)} = \frac{d[\int K f(x) dx - \int f(x) \log f(x) dx]}{d[f(x)]} = 0$$

this is equal to:

$$\begin{aligned} \frac{d[\int K f(x) dx - \int f(x) \log f(x) dx]}{dx} \times dx d[f(x)] &= 0 \\ [K f(x) - f(x) \log f(x)] \times \frac{1}{f'(x)} &= 0 \end{aligned}$$

which is :

$$Kf'(x) - [f'(x)\log f(x) + f(x)\frac{1}{f(x)}f'(x)] = 0$$

and:

$$K - \log f(x) - 1 = 0$$

which means the argmax of ELBO can be find at:

$$E_{-j}[\log p(z_j|z_{-j}, x)] - \log q(z_j) - 1 = 0$$

$$\log \frac{e^{E_{-j}[\log p(z_j|z_{-j}, x)]}}{q(z_j)} = \log_e e$$

Or you can simply view it as $y - x - 1 = 0$. This lead to the conclusion:

$$q^*(z_j) \propto \exp \{E_{-j}[\log p(z_j|z_{-j}, x)]\}$$

since $p(z_j|z_{-j}, x) = \frac{p(z_j, z_{-j}, x)}{p(z_{-j}, x)}$, and $p(z_{-j})$ does not depend on z_j we can equivalently write:

$$q^*(z_j) \propto \exp \{E_{-j}[\log p(z_j, z_{-j}, x)]\}$$

Algorithm 1 Coordinate Ascent Variational Inference

Input: A model $p(x, z)$, a dataset x
Output: A variational density $q(z) = \prod_{j=1}^m q_j(z_j)$
init: variational factors $q_j(z_j)$
while *ELBO has not converged* **do**
 for all $j \in \{1, \dots, m\}$ **do**
 set $q_j(z_j) \propto \exp \{E_{-j}[\log p(z_j, z_{-j}, x)]\}$
 end for
 compute $ELBO = E_q[\log p(z, x)] - E_q[\log q(z)]$
end while

Note that there is generally no guarantee of convexity of ELBO, this coordinate ascent procedure converges to a local maximum.

We can find this method is closely related to Gibbs sampling. Actually Gibbs sampling is a very classical approximate inference method. In variational inference, we take the expected log and set each variables variational factor iteratively.

5 Latent Dirichlet Allocation

LDA is a conditionally conjugate topic model comparing to PLSA. It treats documents as containing multiple topics, where a topic is a distribution over words in vocabulary.

Before we introduce the relationship with variational inference, let's go through some basic knowledge.

5.1 Conjugate Distributions

In Bayesian probability theory, if the posterior distribution $p(\theta|x)$ is in same probability distribution family as the prior distribution $p(\theta)$, the prior and posterior are then called **conjugate distributions**. And the prior is called **conjugate prior** for the likelihood function.

5.1.1 Beta-Binomial Distribution

Beta distribution is a conjugate prior for binomial distribution. For better understanding, Let's consider a simple problem here.

Question 1 For $X_1, \dots, X_n \sim Uniform(0, 1)$, what is the distribution of K th largest number?

Answer 1 Let's assume the K th largest number locates in $[x, x + \Delta x]$. So there are $K-1$ numbers locate in $(0, x)$, and $n-K$ numbers locate in $(x + \Delta x, 1)$. Describe this in math:

$$P(x \leq X_k \leq x + \Delta x) = n \binom{n-1}{k-1} x^{k-1} (1-x-\Delta x)^{n-k} \Delta x$$

Convert to PDF, set $\alpha = k, \beta = n - k + 1$:

$$f(x) = \lim_{\Delta x \rightarrow 0} \frac{P(x \leq X_k \leq x + \Delta x)}{\Delta x} = \frac{\Gamma(\alpha + \beta)}{\Gamma(\alpha)\Gamma(\beta)} x^{\alpha-1} (1-x)^{\beta-1}$$

$$f(x) = \frac{x^{\alpha-1} (1-x)^{\beta-1}}{B(\alpha, \beta)}, B(\alpha, \beta) = \frac{\Gamma(\alpha)\Gamma(\beta)}{\Gamma(\alpha + \beta)}$$

This can also be denoted by:

$$X \sim Beta(\alpha, \beta)$$

Question 2 For $X_1, \dots, X_n \sim Uniform(0, 1)$, what is the distribution of K th largest number?

Given the knowledge: $Y_1, \dots, Y_m \sim Uniform(0, 1)$, there are m_1 Y_i smaller than X_k , m_2 Y_i larger than X_k .

Answer 2 Describe the question in math, we have prior $X_k \sim Beta(\alpha, \beta)$, likelihood $m_1 \sim B(m, X_k)$, the posterior PDF now can be written as:

$$X \sim Beta(\alpha + m_1, \beta + m_2)$$

Actually you can treat X, Y as $X_n, X'_m \sim Uniform(0, 1)$. This question is the same as the previous one.

5.1.2 Dirichlet-Multinomial Distribution

Dirichlet-multinomial distribution is also called as ploya distribution. It is just a multivariate extension of beta-binomial distribution. Dirichlet is the high dimensional Beta distribution, like binomial to multivariate.

$$f(x) = \frac{1}{B(\boldsymbol{\alpha})} \prod_{i=1}^K x_i^{\alpha_i - 1}$$

where

$$B(\boldsymbol{\alpha}) = \frac{\prod_{i=1}^K \Gamma(\alpha_i)}{\Gamma(\sum_{i=1}^K \alpha_i)}$$

And we also have the same property:

$$Prior_{Dir} + Likelihood_{Multi} = Posterior_{Dir}$$

$$Dir(p|\boldsymbol{\alpha}) + Multi(\mathbf{m}) = Dir(p|\boldsymbol{\alpha} + \mathbf{m})$$

5.2 Steps of LDA

Let K be a specific number of topics and V the size of vocabulary. LDA defines the following generative process:

1. For each topic in $i = 1, \dots, K$:
 - (a) draw a distribution over words $\phi_i \sim Dir_V(\beta)$
2. For each document in $j=1, \dots, M$:
 - (a) draw a vector of topic proportions $\theta_j \sim Dir_K(\boldsymbol{\alpha})$
 - (b) For each word in $t = 1, \dots, N_j$:
 - i. draw a topic assignment $z_{j,t}^k \sim Mult(\theta_j)$
 - ii. draw a word $w_{j,t}^v \sim Mult(\phi_{z_{j,t}^k})$

Here β is a fixed parameter of dirichlet prior on topics ϕ with respect to V . and $\boldsymbol{\alpha}$ are fixed parameters of dirichlet prior of topics on each document.

Which can also be denoted as:

$$p(\theta, z, \phi, w|\boldsymbol{\alpha}, \beta) = \prod_{i=1}^K p(\phi_i|\beta) \prod_{j=1}^M \left(p(\theta_j|\boldsymbol{\alpha}) \prod_{t=1}^{N_j} p(z_{j,t}|\theta_j) p(w_{j,t}|\phi_{z_{j,t}}) \right)$$

5.3 Variational Inference for LDA

Remember Section 4.1, we mentioned that a typical variational inference problem is:

$$p(z|x, \alpha) = \frac{p(z, x|\alpha)}{\int_z p(z, x|\alpha) dz}$$

where α is fixed parameter and z is hidden variable. if we find $\int_z p(z, x|\alpha) dz$ intractable, we just create a new distribution $q(z|v)$ to approximate the posterior.

Here for LDA, we have:

$$p(\theta, \phi, z|w, \alpha, \beta) = \frac{p(\theta, \phi, z, w|\alpha, \beta)}{\int_\theta \int_\phi \sum_z p(\theta, \phi, z, w|\alpha, \beta) d_\phi d_\theta}$$

where α, β are some fixed parameters, and θ, ϕ, z are hidden parameters.

The topic assignments z and their prior distribution θ are in conjugate relationship, which is good. But the introduction of ϕ is in coupling with z which makes the posterior intractable.

So we can build a tractable distribution $q(\phi, \theta, z|\lambda, \gamma, \pi)$, with variational variable λ, γ, π

$$q(\phi, \theta, z|\lambda, \gamma, \pi) = \prod_{i=1}^K q(\phi_i|\lambda_i) \prod_{j=1}^M \left(q(\theta_j|\gamma_j) \prod_{t=1}^{N_j} q(z_{j,t}|\pi_{j,t}) \right)$$

And we can measure the KL divergence between $p(\cdot)$ and $q(\cdot)$:

$$\mathcal{L} = \arg \min_{\lambda, \gamma, \pi} KL(q(\phi, \theta, z|\lambda, \gamma, \pi) || p(\theta, \phi, z|w, \alpha, \beta))$$

Since LDA is a conditionally conjugate model. We can directly identify the family of each variable:

$$q(\phi|\lambda) \sim Dir_v(\lambda) \sim Dir_v \left(\beta + \sum_{j=1}^M \sum_{t=1}^N z_{j,t} w_{j,t} \right)$$

$$q(\theta|\gamma) \sim Dir_k(\gamma) \sim Dir_k(\alpha + \sum_{t=1}^N z_{j,t})$$

$$q(z|\pi) \sim Multi(\pi) \propto \exp(\log \theta_{j,k} + \log \phi_{k, w_{dn}})$$

This is just the same question as section 4.3.2. We can calculate the ELBO and update the variational variable until ELBO converge:

E-step: for each document j , each word t , compute π, γ until converge

M-step: using $q(\cdot)$ re-estimate λ